

```
In [1]: def ista_solve_hot( A, d, la_array ):
# ista_solve_hot: Iterative soft-thresholding for multiple values of
# lambda with hot start for each case - the converged value for the previous
# value of lambda is used as an initial condition for the current lambda.
# this function solves the minimization problem
# Minimize |Ax-d|_2^2 + lambda*|x|_1 (Lasso regression)
# using iterative soft-thresholding.
max_iter = 10**4
tol = 10**(-3)
tau = 1/np.linalg.norm(A,2)**2
n = A.shape[1]
w = np.zeros((n,1))
num_lam = len(la_array)
X = np.zeros((n, num_lam))
for i, each_lambda in enumerate(la_array):
    for j in range(max_iter):
        z = w - tau*(A.T*(A*w-d))
        w_old = w
        w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
        X[:, i:i+1] = w
        if np.linalg.norm(w - w_old) < tol:
            break
    return X
```

1a)

```
In [2]: import numpy as np
from scipy.io import loadmat
import matplotlib.pyplot as plt
import pickle

X = loadmat("BreastCancer.mat")['X']
y = loadmat("BreastCancer.mat")['y']

X_100 = X[:100,:]
y_100 = y[:100,:]

lam = np.logspace(-8, np.log10(20), 100)


w_star = ista_solve_hot(X, y, lam)

print((X_100@w_star-y_100).shape)
coord1vals = []
coord2vals = []
for c in range(100):
    temp1 = np.linalg.norm(w_star[:,[c]], ord=1)
    coord1vals.append(temp1)
    temp2 = ((np.sum((X_100@w_star[:,[c]]-y_100)**2))**0.5)
    #temp2 = (((np.sum((X_100@w_star-y_100)**2))**2))**0.5
    coord2vals.append(temp2)

plt.plot(coord1vals, coord2vals, 'r')
plt.xlabel("Norm of w_star")
plt.ylabel("Norm of residual")
```

(100, 100)

Out[2]: Text(0, 0.5, 'Norm of residual')



The lower l1 norm represents a lower lambda value. Small values of lambda increase the squared error. Thus we obtain such a graph where for high norm of w^* we have a lower norm of residual.

1b)

```

In [3]: def sparsity(w_arr):
    nonzeroCount = 0
    for i in w_arr:
        if i>10**-6:
            nonzeroCount += 1
    return nonzeroCount


nonzero_values = []
err = []

for i in range(100):
    nonzero_values.append(sparsity(w_star[:,i]))
    y_hat2 = np.sign(X_100@w_star[:,i:i+1])
    if np.all((y_hat2==0)):
        err.append(1)
        continue
    #err.append(error_rate(np.sign(X_100@w[:,i:i+1], y_train)))
    err.append(np.sum(np.abs(y_hat2-y_100))/2/100)

plt.plot(nonzero_values, err, 'r')
plt.xlabel("Number of Nonzero Coefficients")
plt.ylabel("Error Rate")

```

Out[3]: Text(0, 0.5, 'Error Rate')



As the number of nonzero coefficients goes up, we get decreasing error rates. Intuitively, this makes perfect sense, as more and more data points are used in the calculation, and few are wasted as zero coefficients. Thus, the error rate tends to 0 at really high nonzero coefficients or low sparsity.

1c)

We repeat the same process with data from rows 101 to beyond.

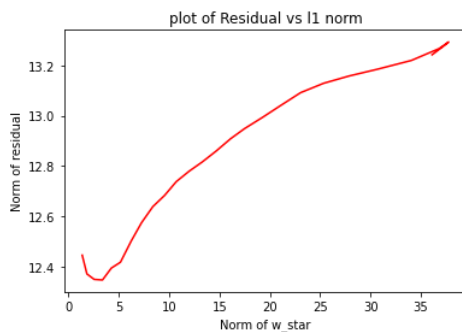
```

In [4]: X_c = X[101:]
    y_c = y[101:]
    lambda_arr = np.logspace(-8,np.log10(20),100)
    w_star = ista_solve_hot(X_100,y_100,lambda_arr)
    coordx = []
    coordy = []
    for c in range(len(w_star[0])):
        w_temp = w_star[:, c:c+1]
        coordx.append(np.linalg.norm(w_temp,ord=1))
        coordy.append(np.linalg.norm(X_c@w_temp-y_c))
    plt.plot(coordx,coordy, 'r')

    plt.xlabel("Norm of w_star")
    plt.ylabel("Norm of residual")
    plt.title("plot of Residual vs l1 norm")

```

Out[4]: Text(0.5, 1.0, 'plot of Residual vs l1 norm')




```

In [5]: coordx = []
        coordy = []
        for i in range(len(w_star[0])):
            w = w_star[:, i:i+1]
            d_hat = np.sign(X_c@w)
            error_vec = [0 if m[0]==m[1] else 1 for m in np.hstack((d_hat, y_c))]
            error_rate = sum(error_vec)/len(error_vec)
            sparsity = 0
            for j in w:
                if j!=0:
                    sparsity+=1
            coordx.append(sparsity)
            coordy.append(error_rate)
        plt.plot(coordx, coordy, 'r')

        plt.xlabel("Sparsity")
        plt.ylabel("Error rate")
        plt.title("Error rate vs sparsity")

```

Out[5]: Text(0.5, 1.0, 'Error rate vs sparsity')



We see a completely different plot for the first plot, whereas in the first case there was an inverse relation between $\|w^*\|_1$ and $\|Aw^* - d\|_2$ in the second case the relation is directly proportional. The error rate wrt sparsity increases a lot initially in this case whereas it does increase as much initially, and instead falls in (b). Thus we conclude that the training set is not a good fit for the data. More features incorporated makes the error rate go up.

In []:

In []:

Q2)

In [1]:

```
def ista_solve_hot( A, d, la_array ):
    # ista_solve_hot: Iterative soft-thresholding for multiple values of
    # lambda with hot start for each case - the converged value for the previous
    # value of lambda is used as an initial condition for the current lambda.
    # this function solves the minimization problem
    # Minimize  $|Ax-d|_2^2 + \lambda|x|_1$  (Lasso regression)
    # using iterative soft-thresholding.
    max_iter = 10**4
    tol = 10**(-3)
    tau = 1/np.linalg.norm(A,2)**2
    n = A.shape[1]
    w = np.zeros((n,1))
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        for j in range(max_iter):
            z = w - tau*(A.T@(A@w-d))
            w_old = w
            w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
            X[:, i:i+1] = w
            if np.linalg.norm(w - w_old) < tol:
                break
    return X

def ridge( A, d, la_array ):
    n = A.shape[1]
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        w = A.T@np.linalg.inv(A@A.T + each_lambda*np.eye(len(A@A.T)))@d
        X[:, i:i+1] = w
    return X
```

```

In [2]: ## Breast Cancer LASSO Exploration
## Prepare workspace
from scipy.io import loadmat
import numpy as np
X = loadmat("BreastCancer.mat")['X']
y = loadmat("BreastCancer.mat")['y']

## 10-fold CV

# each row of setindices denotes the starting and ending index for one
# partition of the data: 5 sets of 30 samples and 5 sets of 29 samples
setindices = [[1,30],[31,60],[61,90],[91,120],[121,150],[151,179],[180,208],[209,237],[238,266],[267,295]]

# each row of holdoutindices denotes the partitions that are held out from
# the training set
holdoutindices = [[1,2],[2,3],[3,4],[4,5],[5,6],[7,8],[9,10],[10,1]]

cases = len(holdoutindices)

# be sure to initiate the quantities you want to measure before looping
# through the various training, validation, and test partitions

avg_error_rate_lasso = 0
avg_sqd_error_ridge = 0
avg_error_rate_ridge = 0
avg_sqd_error_lasso = 0
lam_vals = np.logspace(-8,np.log10(20),100)

# Loop over various cases
for j in range(cases):
    print("Iteration: ", j+1)
    # row indices of first validation set
    v1_ind = np.arange(setindices[holdoutindices[j][0]-1][0]-1,setindices[holdoutindices[j][0]-1][1])

    # row indices of second validation set
    v2_ind = np.arange(setindices[holdoutindices[j][1]-1][0]-1,setindices[holdoutindices[j][1]-1][1])

    # row indices of training set
    trn_ind = list(set(range(295))-set(v1_ind)-set(v2_ind))

    # define matrix of features and labels corresponding to first
    # validation set
    Av1 = X[v1_ind,:]
    bv1 = y[v1_ind]

    # define matrix of features and labels corresponding to second
    # validation set
    Av2 = X[v2_ind,:]
    bv2 = y[v2_ind]

    # define matrix of features and labels corresponding to the
    # training set
    At = X[trn_ind,:]
    bt = y[trn_ind]

    print(len(v1_ind), len(v2_ind), len(trn_ind))

# Use training data to learn classifier weights
wmat_lasso = ista_solve_hot(At,bt,lam_vals)
wmat_ridge = ridge(At,bt,lam_vals)
lam_lasso = None
w_opt_lasso = None
min_error_lasso = -1
lam_ridge = None
w_opt_ridge = None
min_error_ridge = -1

# Find best lambda value using the first validation set, then evaluate everything from there
for i in range(len(wmat_lasso[0])):
    w_lasso = wmat_lasso[:, i:i+1]
    d_hat_lasso = np.sign(Av1@w_lasso)
    error_vec_lasso = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_lasso, bv1))]
    error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
    if min_error_lasso == -1 or error_rate_lasso < min_error_lasso:
        min_error_lasso = error_rate_lasso
        lam_lasso = lam_vals[i]
        w_opt_lasso = w_lasso
    w_ridge = wmat_ridge[:, i:i+1]
    d_hat_ridge = np.sign(Av1@w_ridge)
    error_vec_ridge = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_ridge, bv1))]
    error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)
    if min_error_ridge == -1 or error_rate_ridge < min_error_ridge:
        min_error_ridge = error_rate_ridge
        lam_ridge = lam_vals[i]
        w_opt_ridge = w_ridge

# perform on second validation set, and accumulate performance metrics over all cases and their respective splittings
print("Best lambda in LASSO: ", lam_lasso)
print("Best lambda in ridge regression: ", lam_ridge)
d_hat_lasso = np.sign(Av2@w_opt_lasso)
error_vec_lasso = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_lasso,bv2))]
error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
avg_error_rate_lasso += error_rate_lasso
squared_error_lasso = np.linalg.norm(d_hat_lasso-bv2)**2
avg_sqd_error_lasso += squared_error_lasso
print("Error rate of lasso: ", error_rate_lasso)
print("Squared error of lasso: ", squared_error_lasso)
print("avg error rate of lasso: ", avg_error_rate_lasso)
d_hat_ridge = np.sign(Av2@w_opt_ridge)
error_vec_ridge = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_ridge,bv2))]
error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)

```

```

    avg_error_rate_ridge += error_rate_ridge
    squared_error_ridge = np.linalg.norm(d_hat_ridge-bv2)**2
    avg_sqd_error_ridge += squared_error_ridge
    print("Error rate of ridge: ", error_rate_ridge)
    print("Squared error of ridge: ", squared_error_ridge)
    print("avg error rate of ridge: ", avg_error_rate_ridge)
    print("\n\n")
avg_error_rate_ridge /= cases
avg_sqd_error_ridge /= cases
avg_error_rate_lasso /= cases
avg_sqd_error_lasso /= cases
print("Average error rate of lasso: ", avg_error_rate_lasso)
print("Average squared error of lasso: ", avg_sqd_error_lasso)
print("Average error rate of ridge: ", avg_error_rate_ridge)
print("Average squared error of ridge: ", avg_sqd_error_ridge)

```

```

Iteration: 1
30 30 235
Best lambda in LASSO: 3.5434932692979846
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.43333333333333335
Squared error of lasso: 51.999999999999999
avg error rate of lasso: 0.43333333333333335
Error rate of ridge: 0.43333333333333335
Squared error of ridge: 51.999999999999999
avg error rate of ridge: 0.43333333333333335

```

```

Iteration: 2
30 30 235
Best lambda in LASSO: 16.109430878355187
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.26666666666666666
Squared error of lasso: 32.000000000000001
avg error rate of lasso: 0.7
Error rate of ridge: 0.23333333333333334
Squared error of ridge: 28.000000000000004
avg error rate of ridge: 0.6666666666666667

```

```

Iteration: 3
30 30 235
Best lambda in LASSO: 3.5434932692979846
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.46666666666666667
Squared error of lasso: 56.0
avg error rate of lasso: 1.1666666666666665
Error rate of ridge: 0.43333333333333335
Squared error of ridge: 51.999999999999999
avg error rate of ridge: 1.1

```

```

Iteration: 4
30 30 235
Best lambda in LASSO: 6.780801107819854
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.13333333333333333
Squared error of lasso: 16.0
avg error rate of lasso: 1.2999999999999998
Error rate of ridge: 0.2
Squared error of ridge: 23.999999999999996
avg error rate of ridge: 1.3

```

```

Iteration: 5
30 29 236
Best lambda in LASSO: 1.4915314679609128
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.3103448275862069
Squared error of lasso: 36.0
avg error rate of lasso: 1.6103448275862067
Error rate of ridge: 0.3103448275862069
Squared error of ridge: 36.0
avg error rate of ridge: 1.610344827586207

```

```

Iteration: 6
29 29 237
Best lambda in LASSO: 16.109430878355187
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.13793103448275862
Squared error of lasso: 16.0
avg error rate of lasso: 1.7482758620689653
Error rate of ridge: 0.1724137931034483
Squared error of ridge: 20.000000000000004
avg error rate of ridge: 1.782758620689655

```

```

Iteration: 7
29 29 237
Best lambda in LASSO: 20.000000000000004
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.5172413793103449
Squared error of lasso: 60.000000000000001
avg error rate of lasso: 2.2655172413793103

```

```
Error rate of ridge: 0.4482758620689655
Squared error of ridge: 51.99999999999999
avg error rate of ridge: 2.231034482758621
```

```
Iteration: 8
29 30 236
Best lambda in LASSO: 3.955000701930882e-07
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.2
Squared error of lasso: 23.999999999999996
avg error rate of lasso: 2.4655172413793105
Error rate of ridge: 0.2
Squared error of ridge: 23.999999999999996
avg error rate of ridge: 2.431034482758621
```

```
Average error rate of lasso: 0.3081896551724138
Average squared error of lasso: 36.5
Average error rate of ridge: 0.3038793103448276
Average squared error of ridge: 36.0
```

In []: