

## Q2)

In [1]:

```
def ista_solve_hot( A, d, la_array ):
    # ista_solve_hot: Iterative soft-thresholding for multiple values of
    # lambda with hot start for each case - the converged value for the previous
    # value of lambda is used as an initial condition for the current lambda.
    # this function solves the minimization problem
    # Minimize  $\|Ax-d\|_2^2 + \lambda \|x\|_1$  (Lasso regression)
    # using iterative soft-thresholding.
    max_iter = 10**4
    tol = 10**(-3)
    tau = 1/np.linalg.norm(A,2)**2
    n = A.shape[1]
    w = np.zeros((n,1))
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        for j in range(max_iter):
            z = w - tau*(A.T*(A@w-d))
            w_old = w
            w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
            X[:, i:i+1] = w
            if np.linalg.norm(w - w_old) < tol:
                break
    return X

def ridge( A, d, la_array ):
    n = A.shape[1]
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        w = A.T@np.linalg.inv(A@A.T + each_lambda*np.eye(len(A@A.T)))@d
        X[:, i:i+1] = w
    return X
```

```

In [2]: ## Breast Cancer LASSO Exploration
## Prepare workspace
from scipy.io import loadmat
import numpy as np
X = loadmat("BreastCancer.mat")['X']
y = loadmat("BreastCancer.mat")['y']

## 10-fold CV

# each row of setindices denotes the starting and ending index for one
# partition of the data: 5 sets of 30 samples and 5 sets of 29 samples
setindices = [[1,30],[31,60],[61,90],[91,120],[121,150],[151,179],[180,208],[209,237],[238,266],[267,295]]

# each row of holdoutindices denotes the partitions that are held out from
# the training set
holdoutindices = [[1,2],[2,3],[3,4],[4,5],[5,6],[7,8],[9,10],[10,1]]

cases = len(holdoutindices)

# be sure to initiate the quantities you want to measure before looping
# through the various training, validation, and test partitions

avg_error_rate_lasso = 0
avg_sqd_error_ridge = 0
avg_error_rate_ridge = 0
avg_sqd_error_lasso = 0
lam_vals = np.logspace(-8,np.log10(20),100)

# Loop over various cases
for j in range(cases):
    print("Iteration: ", j+1)
    # row indices of first validation set
    v1_ind = np.arange(setindices[holdoutindices[j][0]-1][0]-1,setindices[holdoutindices[j][0]-1][1])

    # row indices of second validation set
    v2_ind = np.arange(setindices[holdoutindices[j][1]-1][0]-1,setindices[holdoutindices[j][1]-1][1])

    # row indices of training set
    trn_ind = list(set(range(295))-set(v1_ind)-set(v2_ind))

    # define matrix of features and labels corresponding to first
    # validation set
    Av1 = X[v1_ind,:]
    bv1 = y[v1_ind]

    # define matrix of features and labels corresponding to second
    # validation set
    Av2 = X[v2_ind,:]
    bv2 = y[v2_ind]

    # define matrix of features and labels corresponding to the
    # training set
    At = X[trn_ind,:]
    bt = y[trn_ind]

    print(len(v1_ind), len(v2_ind), len(trn_ind))

# Use training data to learn classifier weights
wmat_lasso = ista_solve_hot(At,bt,lam_vals)
wmat_ridge = ridge(At,bt,lam_vals)
lam_lasso = None
w_opt_lasso = None
min_error_lasso = -1
lam_ridge = None
w_opt_ridge = None
min_error_ridge = -1

# Find best lambda value using the first validation set, then evaluate everything from there
for i in range(len(wmat_lasso[0])):
    w_lasso = wmat_lasso[:, i:i+1]
    d_hat_lasso = np.sign(Av1@w_lasso)
    error_vec_lasso = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_lasso, bv1))]
    error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
    if min_error_lasso == -1 or error_rate_lasso < min_error_lasso:
        min_error_lasso = error_rate_lasso
        lam_lasso = lam_vals[i]
        w_opt_lasso = w_lasso
    w_ridge = wmat_ridge[:, i:i+1]
    d_hat_ridge = np.sign(Av1@w_ridge)
    error_vec_ridge = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_ridge, bv1))]
    error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)
    if min_error_ridge == -1 or error_rate_ridge < min_error_ridge:
        min_error_ridge = error_rate_ridge
        lam_ridge = lam_vals[i]
        w_opt_ridge = w_ridge

# perform on second validation set, and accumulate performance metrics over all cases and their respective splittings
print("Best lambda in LASSO: ", lam_lasso)
print("Best lambda in ridge regression: ", lam_ridge)
d_hat_lasso = np.sign(Av2@w_opt_lasso)
error_vec_lasso = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_lasso,bv2))]
error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
avg_error_rate_lasso += error_rate_lasso
squared_error_lasso = np.linalg.norm(d_hat_lasso-bv2)**2
avg_sqd_error_lasso += squared_error_lasso
print("Error rate of lasso: ", error_rate_lasso)
print("Squared error of lasso: ", squared_error_lasso)
print("avg error rate of lasso: ", avg_error_rate_lasso)
d_hat_ridge = np.sign(Av2@w_opt_ridge)
error_vec_ridge = [0 if m[1]==m[0] else 1 for m in np.hstack((d_hat_ridge,bv2))]
error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)

```

```

    avg_error_rate_ride += error_rate_ride
    squared_error_ride = np.linalg.norm(d_hat_ride-bv2)**2
    avg_sqd_error_ride += squared_error_ride
    print("Error rate of ridge: ", error_rate_ride)
    print("Squared error of ridge: ", squared_error_ride)
    print("avg error rate of ridge: ", avg_error_rate_ride)
    print("\n\n")
avg_error_rate_ride /= cases
avg_sqd_error_ride /= cases
avg_error_rate_lasso /= cases
avg_sqd_error_lasso /= cases
print("Average error rate of lasso: ", avg_error_rate_lasso)
print("Average squared error of lasso: ", avg_sqd_error_lasso)
print("Average error rate of ridge: ", avg_error_rate_ride)
print("Average squared error of ridge: ", avg_sqd_error_ride)

```

```

Iteration: 1
30 30 235
Best lambda in LASSO: 3.5434932692979846
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.43333333333333335
Squared error of lasso: 51.999999999999999
avg error rate of lasso: 0.43333333333333335
Error rate of ridge: 0.43333333333333335
Squared error of ridge: 51.999999999999999
avg error rate of ridge: 0.43333333333333335

```

```

Iteration: 2
30 30 235
Best lambda in LASSO: 16.109430878355187
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.26666666666666666
Squared error of lasso: 32.000000000000001
avg error rate of lasso: 0.7
Error rate of ridge: 0.23333333333333334
Squared error of ridge: 28.000000000000004
avg error rate of ridge: 0.6666666666666667

```

```

Iteration: 3
30 30 235
Best lambda in LASSO: 3.5434932692979846
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.46666666666666667
Squared error of lasso: 56.0
avg error rate of lasso: 1.1666666666666665
Error rate of ridge: 0.43333333333333335
Squared error of ridge: 51.999999999999999
avg error rate of ridge: 1.1

```

```

Iteration: 4
30 30 235
Best lambda in LASSO: 6.780801107819854
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.13333333333333333
Squared error of lasso: 16.0
avg error rate of lasso: 1.2999999999999998
Error rate of ridge: 0.2
Squared error of ridge: 23.999999999999996
avg error rate of ridge: 1.3

```

```

Iteration: 5
30 29 236
Best lambda in LASSO: 1.4915314679609128
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.3103448275862069
Squared error of lasso: 36.0
avg error rate of lasso: 1.6103448275862067
Error rate of ridge: 0.3103448275862069
Squared error of ridge: 36.0
avg error rate of ridge: 1.610344827586207

```

```

Iteration: 6
29 29 237
Best lambda in LASSO: 16.109430878355187
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.13793103448275862
Squared error of lasso: 16.0
avg error rate of lasso: 1.7482758620689653
Error rate of ridge: 0.1724137931034483
Squared error of ridge: 20.000000000000004
avg error rate of ridge: 1.782758620689655

```

```

Iteration: 7
29 29 237
Best lambda in LASSO: 20.000000000000004
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.5172413793103449
Squared error of lasso: 60.000000000000001
avg error rate of lasso: 2.2655172413793103

```

```
Error rate of ridge: 0.4482758620689655
Squared error of ridge: 51.99999999999999
avg error rate of ridge: 2.231034482758621
```

```
Iteration: 8
29 30 236
Best lambda in LASSO: 3.955000701930882e-07
Best lambda in ridge regression: 1e-08
Error rate of lasso: 0.2
Squared error of lasso: 23.999999999999996
avg error rate of lasso: 2.4655172413793105
Error rate of ridge: 0.2
Squared error of ridge: 23.999999999999996
avg error rate of ridge: 2.431034482758621
```

```
Average error rate of lasso: 0.3081896551724138
Average squared error of lasso: 36.5
Average error rate of ridge: 0.3038793103448276
Average squared error of ridge: 36.0
```

In [ ]: