

ista_solve_hot

April 13, 2021

1 Q1

```
[172]: import numpy as np
import scipy.io as sp
import matplotlib.pyplot as plt
```

```
[173]: data = sp.loadmat("BreastCancer.mat")
```

```
[174]: def ista_solve_hot( A, d, la_array ):
    # ista_solve_hot: Iterative soft-thresholding for multiple values of
    # lambda with hot start for each case - the converged value for the previous
    # value of lambda is used as an initial condition for the current lambda.
    # this function solves the minimization problem
    # Minimize ||Ax-d||_2^2 + lambda*||x||_1 (Lasso regression)
    # using iterative soft-thresholding.
    max_iter = 10**4
    tol = 10**(-3)
    tau = 1/np.linalg.norm(A,2)**2
    n = A.shape[1]
    w = np.zeros((n,1))
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        for j in range(max_iter):
            z = w - tau*(A.T@(A@w-d))
            w_old = w
            w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
            X[:, i:i+1] = w
            if np.linalg.norm(w - w_old) < tol:
                break
    return X
```

1.1 (a)

```
[175]: A = data['X'][:100]
d = data['y'][:100]
lambda_arr = np.logspace(-8,np.log10(20),100)

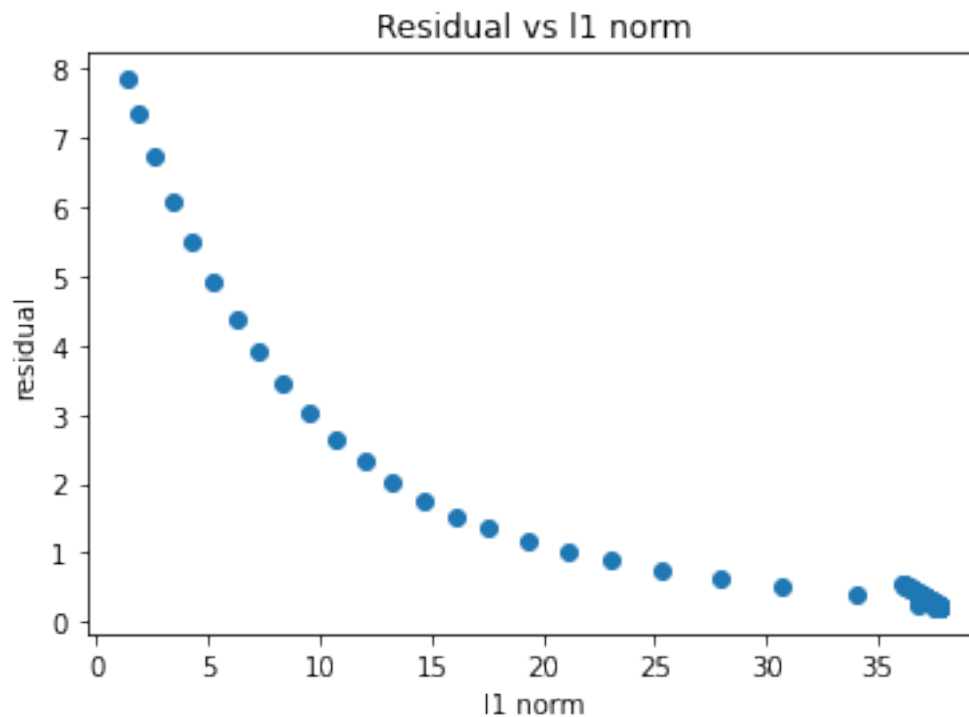
wmat = ista_solve_hot(A,d,lambda_arr)
```

```
[176]: plot_x = []
plot_y = []

for i in range(len(wmat[0])):
    w = wmat[:, i:i+1]
    plot_x.append(np.linalg.norm(w,ord=1))
    plot_y.append(np.linalg.norm(A@w-d))

plt.title("Residual vs l1 norm")
plt.xlabel("l1 norm")
plt.ylabel("residual")
plt.scatter(plot_x,plot_y)
```

```
[176]: <matplotlib.collections.PathCollection at 0x7f51dc63f390>
```



1.2 The residual and l1 norm seem to have an inverse relationship


1.3 (b)

```
[177]: plot_x = []
plot_y = []

for i in range(len(wmat[0])):
    w = wmat[:, i:i+1]
    d_hat = np.sign(A@w)
    error_vec = [0 if m[0]==m[1] else 1 for m in np.hstack((d_hat, d))]
    error_rate = sum(error_vec)/len(error_vec)
    sparsity = 0
    for j in w:
        if j!=0:
            sparsity+=1
    plot_x.append(sparsity)
    plot_y.append(error_rate)

plt.title("Error rate vs sparsity")
plt.xlabel("sparsity")
plt.ylabel("Error rate")
plt.scatter(plot_x,plot_y)
```

[177]: <matplotlib.collections.PathCollection at 0x7f51dc5afd50>



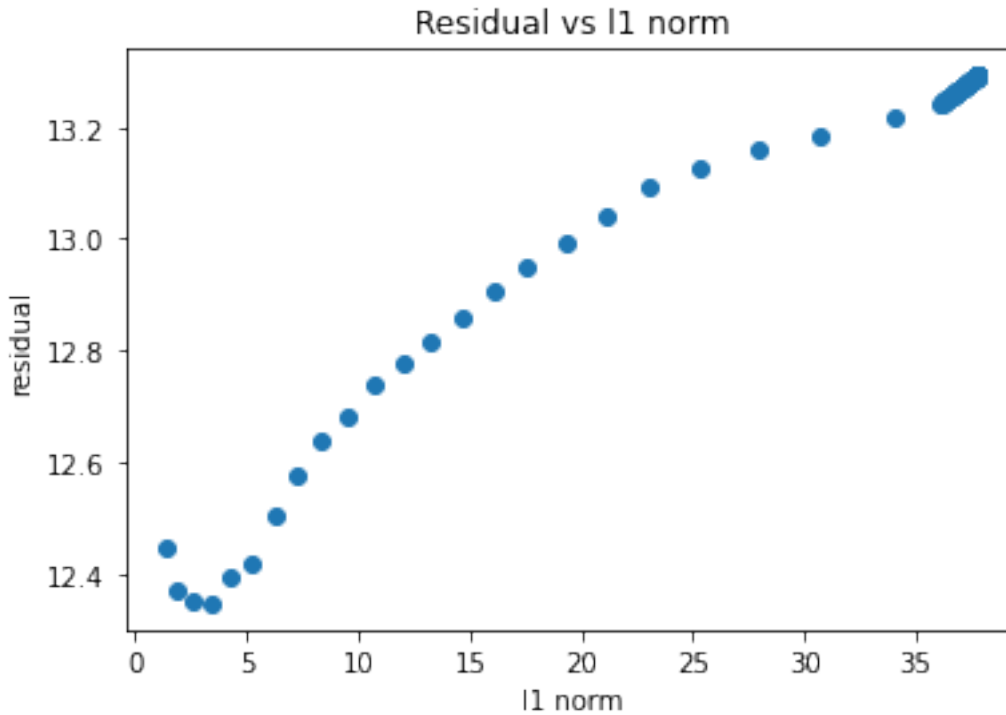
- 1.4 As the sparsity goes up, the error rate goes down. Initially, increasing the sparsity even by a little bit makes the error rate drop drastically. The drop in error rate becomes insignificant beyond a particular sparsity telling that a good sparse solution exists which focuses on only a few significant features

2 (c)

```
[178]: A_test = data['X'][101:]  
       d_test = data['y'][101:]
```

```
[179]: plot_x = []  
       plot_y = []  
  
       for i in range(len(wmat[0])):  
           w = wmat[:, i:i+1]  
           plot_x.append(np.linalg.norm(w,ord=1))  
           plot_y.append(np.linalg.norm(A_test@w-d_test))  
  
       plt.title("Residual vs l1 norm")  
       plt.xlabel("l1 norm")  
       plt.ylabel("residual")  
       plt.scatter(plot_x,plot_y)
```

```
[179]: <matplotlib.collections.PathCollection at 0x7f51dc5a2e90>
```



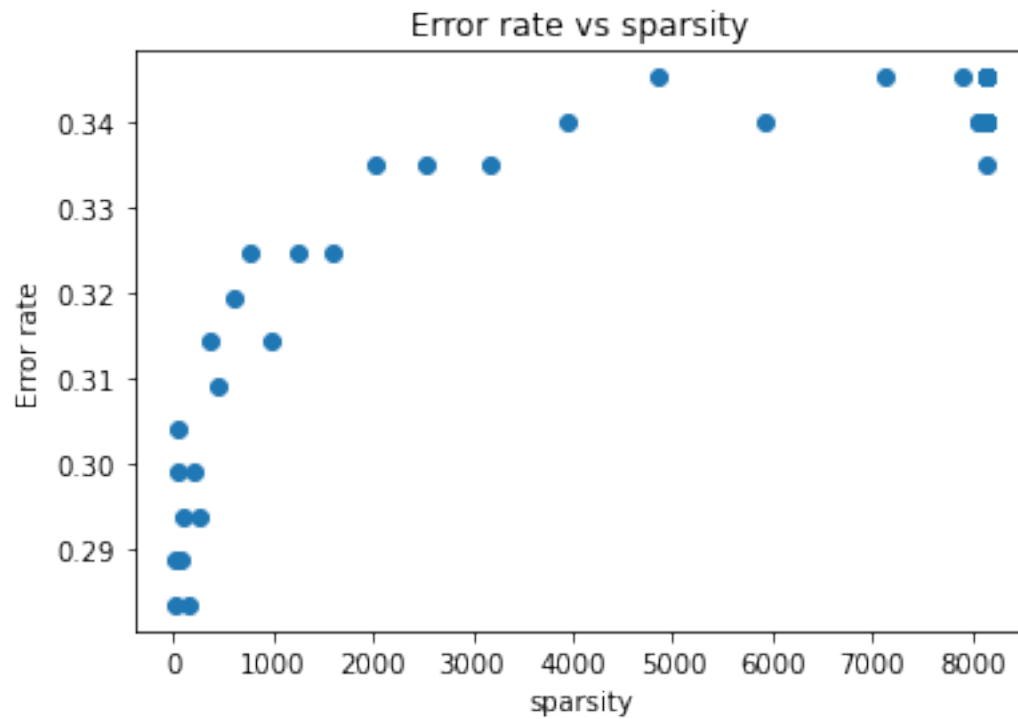
2.1 There seems to be a linear relationship between residual and l1 norm, contrary to the previous graph from our training data

```
[180]: plot_x = []
plot_y = []

for i in range(len(wmat[0])):
    w = wmat[:, i:i+1]
    d_hat = np.sign(A_test@w)
    error_vec = [0 if m[0]==m[1] else 1 for m in np.hstack((d_hat, d_test))]
    error_rate = sum(error_vec)/len(error_vec)
    sparsity = 0
    for j in w:
        if j!=0:
            sparsity+=1
    plot_x.append(sparsity)
    plot_y.append(error_rate)

plt.title("Error rate vs sparsity")
plt.xlabel("sparsity")
plt.ylabel("Error rate")
plt.scatter(plot_x,plot_y)
```

[180]: <matplotlib.collections.PathCollection at 0x7f51dc513ad0>



2.2 Here, the error rate seems to increase by a lot when increasing sparsity even by a bit. Our model does not seem to be a good fit for the data because taking more features into consideration makes the error rate go up

[]:

Q2_starter

April 13, 2021

1 Q2

```
[26]: def ista_solve_hot( A, d, la_array ):  
    # ista_solve_hot: Iterative soft-thresholding for multiple values of  
    # lambda with hot start for each case - the converged value for the previous  
    # value of lambda is used as an initial condition for the current lambda.  
    # this function solves the minimization problem  
    # Minimize  $\|Ax-d\|_2^2 + \lambda\|x\|_1$  (Lasso regression)  
    # using iterative soft-thresholding.  
    max_iter = 10**4  
    tol = 10**(-3)  
    tau = 1/np.linalg.norm(A,2)**2  
    n = A.shape[1]  
    w = np.zeros((n,1))  
    num_lam = len(la_array)  
    X = np.zeros((n, num_lam))  
    for i, each_lambda in enumerate(la_array):  
        for j in range(max_iter):  
            z = w - tau*(A.T*(A@w-d))  
            w_old = w  
            w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)  
            X[:, i:i+1] = w  
            if np.linalg.norm(w - w_old) < tol:  
                break  
    return X
```

```
[27]: def ridge( A, d, la_array ):  
    n = A.shape[1]  
    num_lam = len(la_array)  
    X = np.zeros((n, num_lam))  
    for i, each_lambda in enumerate(la_array):  
        w = A.T*np.linalg.inv(A@A.T + each_lambda*np.eye(len(A@A.T)))@d  
        X[:, i:i+1] = w  
    return X
```

```

[28]: ## Breast Cancer LASSO Exploration
      ## Prepare workspace
      from scipy.io import loadmat
      import numpy as np
      X = loadmat("BreastCancer.mat")['X']
      y = loadmat("BreastCancer.mat")['y']

      ## 10-fold CV

      # each row of setindices denotes the starting an ending index for one
      # partition of the data: 5 sets of 30 samples and 5 sets of 29 samples
      setindices =         
      → [[1,30],[31,60],[61,90],[91,120],[121,150],[151,179],[180,208],[209,237],[238,266],[267,295]]

      # each row of holdoutindices denotes the partitions that are held out from
      # the training set
      holdoutindices = [[1,2],[2,3],[3,4],[4,5],[5,6],[7,8],[9,10],[10,1]]

      cases = len(holdoutindices)

      # be sure to initiate the quantities you want to measure before looping
      # through the various training, validation, and test partitions
      avg_sqd_error_lasso = 0
      avg_sqd_error_ridge = 0
      avg_error_rate_ridge = 0
      avg_error_rate_lasso = 0
      lam_vals = np.logspace(-8,np.log10(20),100)

      # Loop over various cases
      for j in range(cases):
          print("Iteration: ", j+1)
          # row indices of first validation set
          v1_ind = np.
          → arange(setindices[holdoutindices[j][0]-1][0]-1,setindices[holdoutindices[j][0]-1][1])

          # row indices of second validation set
          v2_ind = np.
          → arange(setindices[holdoutindices[j][1]-1][0]-1,setindices[holdoutindices[j][1]-1][1])

          # row indices of training set
          trn_ind = list(set(range(295))-set(v1_ind)-set(v2_ind))

          # define matrix of features and labels corresponding to first
          # validation set
          Av1 = X[v1_ind,:]
          bv1 = y[v1_ind]

```



```

# define matrix of features and labels corresponding to second
# validation set
Av2 = X[v2_ind,:]
bv2 = y[v2_ind]

# define matrix of features and labels corresponding to the
# training set
At = X[trn_ind,:]
bt = y[trn_ind]

print("Set sizes")
print("v1: ", len(v1_ind), "v2: ", len(v2_ind), "training: ", len(trn_ind))
# Use training data to learn classifier
wmat_lasso = ista_solve_hot(At,bt,lam_vals)
wmat_ridge = ridge(At,bt,lam_vals)
lam_lasso = None
w_opt_lasso = None
min_error_lasso = -1
lam_ridge = None
w_opt_ridge= None
min_error_ridge = -1
# Find best lambda value using first validation set
for i in range(len(wmat_lasso[0])):
    w_lasso = wmat_lasso[:, i:i+1]
    d_hat_lasso = np.sign(Av1@w_lasso)
    error_vec_lasso = [0 if m[0]==m[1] else 1 for m in np.
→hstack((d_hat_lasso, bv1))]
    error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
    if min_error_lasso == -1 or error_rate_lasso<min_error_lasso:
        min_error_lasso = error_rate_lasso
        lam_lasso = lam_vals[i]
        w_opt_lasso = w_lasso

    w_ridge = wmat_ridge[:, i:i+1]
    d_hat_ridge = np.sign(Av1@w_ridge)
    error_vec_ridge = [0 if m[0]==m[1] else 1 for m in np.
→hstack((d_hat_ridge, bv1))]
    error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)
    if min_error_ridge == -1 or error_rate_ridge<min_error_ridge:
        min_error_ridge = error_rate_ridge
        lam_ridge = lam_vals[i]
        w_opt_ridge = w_ridge

# Evaluate performance on second validation set
# and accumulate performance metrics over all cases partitions
print("Best lambda for LASSO: ", lam_lasso)
print("Best lambda for ridge regression: ", lam_ridge)

```

```

d_hat_lasso = np.sign(Av2@w_opt_lasso)
error_vec_lasso = [0 if m[0]==m[1] else 1 for m in np.hstack((d_hat_lasso,
↳bv2))]
error_rate_lasso = sum(error_vec_lasso)/len(error_vec_lasso)
avg_error_rate_lasso += error_rate_lasso
squared_error_lasso = np.linalg.norm(d_hat_lasso-bv2)**2
avg_sqd_error_lasso += squared_error_lasso
print("Error rate lasso: ", error_rate_lasso)
print("Squared error lasso: ", squared_error_lasso)

d_hat_ridge = np.sign(Av2@w_opt_ridge)
error_vec_ridge = [0 if m[0]==m[1] else 1 for m in np.hstack((d_hat_ridge,
↳bv2))]
error_rate_ridge = sum(error_vec_ridge)/len(error_vec_ridge)
avg_error_rate_ridge += error_rate_ridge
squared_error_ridge = np.linalg.norm(d_hat_ridge-bv2)**2
avg_sqd_error_ridge += squared_error_ridge
print("Error rate ridge: ", error_rate_ridge)
print("Squared error ridge: ", squared_error_ridge)

print("\n")

avg_error_rate_ridge /= cases
avg_sqd_error_ridge /= cases
avg_error_rate_lasso /= cases
avg_sqd_error_lasso /= cases

print("#####")
print("Average error rate lasso: ", avg_error_rate_lasso)
print("Average squared error lasso: ", avg_sqd_error_lasso)
print("Average error rate ridge: ", avg_error_rate_ridge)
print("Average squared error ridge: ", avg_sqd_error_ridge)

```

```

Iteration: 1
Set sizes
v1: 30 v2: 30 training: 235
Best lambda for LASSO: 3.5434932692979846
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.43333333333333335
Squared error lasso: 51.99999999999999
Error rate ridge: 0.43333333333333335
Squared error ridge: 51.99999999999999

```

```

Iteration: 2
Set sizes

```

v1: 30 v2: 30 training: 235
Best lambda for LASSO: 16.109430878355187
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.2666666666666666
Squared error lasso: 32.00000000000001
Error rate ridge: 0.23333333333333334
Squared error ridge: 28.000000000000004

Iteration: 3
Set sizes
v1: 30 v2: 30 training: 235
Best lambda for LASSO: 3.5434932692979846
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.4666666666666667
Squared error lasso: 56.0
Error rate ridge: 0.43333333333333335
Squared error ridge: 51.99999999999999

Iteration: 4
Set sizes
v1: 30 v2: 30 training: 235
Best lambda for LASSO: 6.780801107819854
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.13333333333333333
Squared error lasso: 16.0
Error rate ridge: 0.2
Squared error ridge: 23.999999999999996

Iteration: 5
Set sizes
v1: 30 v2: 29 training: 236
Best lambda for LASSO: 1.4915314679609128
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.3103448275862069
Squared error lasso: 36.0
Error rate ridge: 0.3103448275862069
Squared error ridge: 36.0

Iteration: 6
Set sizes
v1: 29 v2: 29 training: 237
Best lambda for LASSO: 16.109430878355187
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.13793103448275862

Squared error lasso: 16.0
Error rate ridge: 0.1724137931034483
Squared error ridge: 20.000000000000004

Iteration: 7
Set sizes
v1: 29 v2: 29 training: 237
Best lambda for LASSO: 20.000000000000004
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.5172413793103449
Squared error lasso: 60.000000000000001
Error rate ridge: 0.4482758620689655
Squared error ridge: 51.99999999999999

Iteration: 8
Set sizes
v1: 29 v2: 30 training: 236
Best lambda for LASSO: 3.955000701930882e-07
Best lambda for ridge regression: 1e-08
Error rate lasso: 0.2
Squared error lasso: 23.999999999999996
Error rate ridge: 0.2
Squared error ridge: 23.999999999999996

Average error rate lasso: 0.3081896551724138
Average squared error lasso: 36.5
Average error rate ridge: 0.3038793103448276
Average squared error ridge: 36.0

[]: